CSE 620: Combinatorial Optimization

Simulated Annealing Micro-Project Report

Cong Hoan Nguyen (c0nguy10@louisville.edu)
and Anesa Hozdic (anesa.hozdic@louisville.edu)
University of Louisville

November 3, 2026

1. Simulated Annealing: Why It Beats Classical Local Search

Simulated annealing (SA) is a stochastic local-search heuristic that explores a solution space by repeatedly moving to a neighboring state, accepting every improving move and accepting a worsening move of size ΔE with probability $\exp(-\Delta E/T)$, where T is a "temperature" that is gradually reduced according to a cooling schedule. Early in the search, a high T makes uphill moves almost always acceptable, so the algorithm behaves like a random walk and explores many basins of attraction; as $T \to 0$ the acceptance probability of worse moves vanishes and the search settles into the best basin it has found, i.e. it exploits. This is precisely what classical methods cannot do: hill climbing and gradient descent are greedy and only accept improving moves, so they are guaranteed to terminate at whatever local optimum (or stationary point) happens to be closest to the starting state, and they require a good initial guess and (for gradient descent) a differentiable objective. Because SA retains a non-zero probability of escaping a local optimum while the temperature is still high, and because a slowly decaying schedule lets it sample the landscape like a Markov chain before freezing, it can reach states that greedy search can never leave — giving it a much better (approaching global, in the limit of logarithmic cooling) approximation of the global optimum, at the price of longer run times and a few extra parameters (T_0, α, schedule type) that must be tuned.

2. Experiment 1 — Exponential Cooling Schedule

Exponential decay cools the temperature as

$$T_k = T_0\,\alpha^k, \tag{1}$$

where T_0 is the starting temperature and $\alpha \in (0, 1)$ is the decay rate. The following table reports the results of running Section 1 of the provided notebook (15-city TSP instance generated with `random_state=254`, SA seeded with `random_state=192`, averaged over 10 runs per configuration, stopping at $T < 10^{-5}$).

Observations. Both parameters control the same underlying quantity — the number of iterations performed, since exponential cooling reaches the 10^{-5} stop temperature after only $k_{\max} = \ln(10^{-5}/T_0)/\ln\alpha$ steps (about 400 steps for $\alpha = 0.95$, but only ∼40 and ∼17 steps for $\alpha = 0.60$ and 0.30). There is a clear positive correlation between α (and, more weakly, T_0) and solution quality, as visualized in Figure 1: the slowest schedule, $T_0 = 4000,\ \alpha = 0.95$, found the best tour (2504.8) while the fastest schedules terminated after a dozen steps and returned tours of 4200–4400, i.e.

Table 1: Exponential decay: average time and best cost over 10 runs.

Starting Temperature	Alpha	Average Time (ms)	Best Cost
4000	0.95	330.38	2504.8
2000	0.95	309.17	3434.4
4000	0.60	33.14	3963.2
2000	0.60	37.58	3963.2
4000	0.30	15.26	4393.6
2000	0.30	14.13	4178.4

roughly 75% worse, in under 16 ms. Runtime grows with α far more than with T_0: at $\alpha = 0.30$ the two starting temperatures are almost indistinguishable, because the temperature has crashed to zero in a handful of iterations no matter where it started, so T_0 only sets the acceptance probability of the first few moves. At $\alpha = 0.95$, raising T_0 from 2000 to 4000 improves the cost from 3434.4 to 2504.8 for essentially the same runtime, because a hotter start keeps uphill moves acceptable longer and lets the search cross higher energy barriers before it freezes.

3. Experiment 2 — Linear Cooling Schedule

Linear decay cools the temperature as

$$T_k = \frac{T_0}{1 + \alpha k}, \tag{2}$$

so the temperature now decreases hyperbolically instead of geometrically. The experiment from Section 2 was repeated with the same instance, seeds and averaging protocol, stopping at $T < 1$ as specified in the notebook.

Table 2: Linear decay: average time and best cost over 10 runs.

Starting Temperature	Alpha	Average Time (ms)	Best Cost
4000	0.95	3631.96	2477.6
2000	0.95	1772.89	3753.6
4000	0.60	5444.18	2477.6
2000	0.60	2798.65	2894.4
4000	0.30	10949.77	2559.2
2000	0.30	5213.73	2477.6

Observations. Linear cooling is dramatically slower than exponential cooling: reaching the stop temperature now takes $k_{\max} \approx (T_0 - 1)/\alpha$ iterations — roughly 4200 steps for $T_0 = 4000,\ \alpha = 0.95$ and over 13 000 steps for $T_0 = 4000,\ \alpha = 0.30$ — so run times jump from tens of milliseconds to 1.8–11 seconds, while the best cost drops from 2504.8 (exp) to 2477.6. Note that the role of α is inverted relative to exponential decay: because α sits in the denominator, a *larger* α cools *faster*, so run time now *decreases* with α and *increases* with T_0. Quality still correlates with the number of iterations performed — every configuration that ran at least ∼4200 steps reached the best cost of 2477.6, while the shortest run (2000/0.95, ∼2100 steps) finished at 3753.6 — but the relationship saturates: 13 000 steps costs 2.5× more time than 4200 steps for no gain (2559.2 vs. 2477.6, within

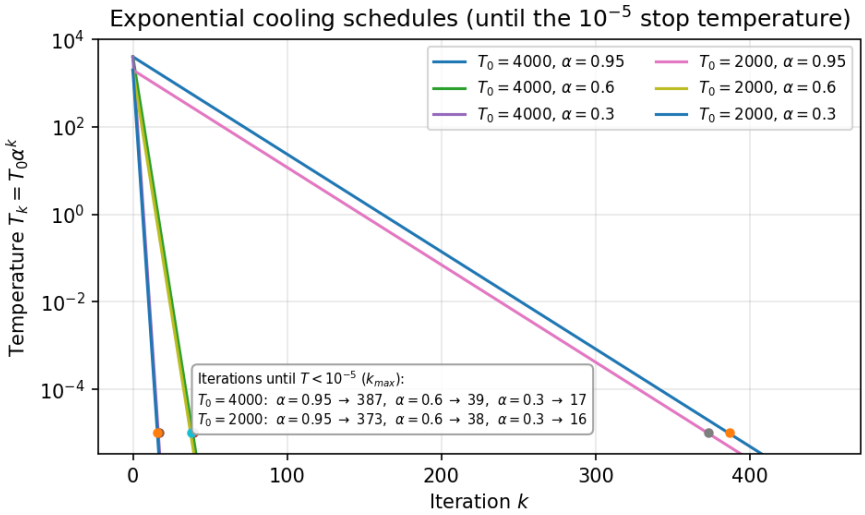


Figure 1: The six exponential schedules tested: a larger α (or larger T_0) keeps the temperature high for many more iterations — each dot marks where a schedule crosses the 10^{-5} stopping threshold, i.e. the number of iterations that configuration performs.

stochastic noise). In other words, with linear decay the budget is controlled by T_0/α, and once enough iterations have been spent for the temperature to become low, extra searching is wasted computation.

4. Experiment 3 — Simulated Annealing vs. Hill Climbing

The best parameter combination from the two tables above is the linear schedule with $T_0 = 4000$ and $\alpha = 0.95$ (best cost 2477.6, and the cheapest of the three configurations that tie at that cost). It was used to solve randomly generated instances of 5, 20 and 30 cities (same `random_state=254` generator) and compared against the greedy hill climber from the notebook on the identical instance.

Table 3: Simulated annealing (best parameters) vs. hill climbing.

	Simulated Annealing		Hill Climbing	
# Cities	Time (ms)	Cost	Time (ms)	Cost
5	550.64	1478.4	0.36	1478.4
20	6511.61	4892.0	20.29	5030.4
30	14606.29	6967.2	44.50	7619.2

Observations. For both algorithms the run time grows with the number of cities — SA because it performs a fixed (temperature-driven) number of iterations whose per-step cost is $O(n^2)$, hill climbing because each greedy sweep evaluates an n-element neighborhood — but they grow on completely different scales: SA needs 0.55 s, 6.5 s and 14.6 s for 5, 20 and 30 cities, while hill climbing needs 0.4 ms, 20 ms and 45 ms, i.e. it is 3–4 orders of magnitude faster throughout. In terms of quality, both algorithms find the same tour for 5 cities (1478.4), because a 5-city neighborhood graph is tiny and the greedy search reaches the global optimum anyway; as the

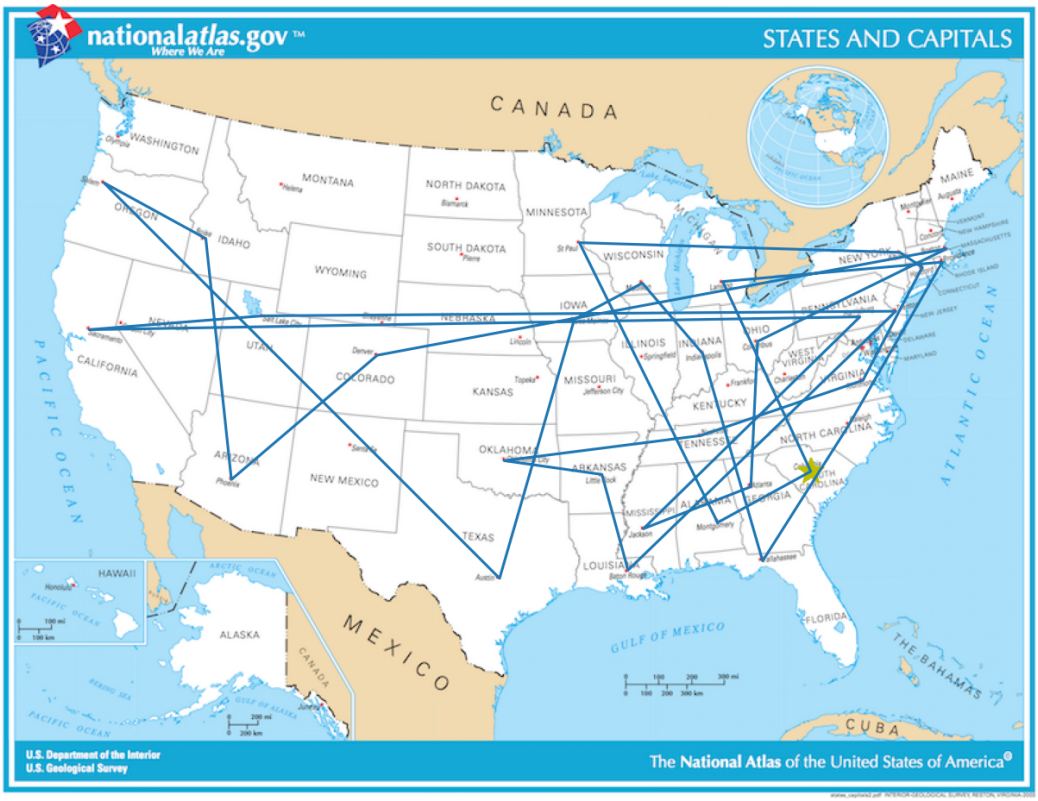


Figure 2: Best 30-city tour found by simulated annealing (linear schedule, $T_0 = 4000$, $\alpha = 0.95$) on the US capitals map.

instance grows, hill climbing freezes in the first local optimum it meets and its cost deteriorates relative to SA: it is 2.8% worse at 20 cities (5030.4 vs. 4892.0) and 9.4% worse at 30 cities (7619.2 vs. 6967.2), and the gap keeps widening with n. For the largest problem sizes, therefore, the best approach is simulated annealing with a slowly cooling schedule, since solution quality — not run time — is what dominates the total routing cost, and exact methods would be out of reach anyway (TSP is NP-hard). For the smallest instance, SA is not worth it: hill climbing returns the same optimal tour about 1500× faster, so the extra stochastic exploration (and its parameter tuning) buys nothing.

5. Concluding Remarks

Parameter and schedule selection make a very significant difference in the effectiveness of simulated annealing. Holding everything else fixed, merely changing α from 0.95 to 0.30 under exponential cooling degraded the tour from 2504.8 to 4393.6 — a 75% worse solution — while changing the cooling schedule from exponential to linear improved the best cost from 2504.8 to 2477.6 but at a 100-fold increase in run time (0.33 s vs. 3.6 s), showing that the schedule controls how the search budget is spent: fast cooling is cheap but premature and greedy-like, slow cooling explores but can waste iterations once the temperature is already low (the 13 000-step linear run gained nothing over the 4 200-step run). Starting temperature matters mainly through the number of iterations it sustains and the acceptance rate it produces early on; it is only consequential when the decay rate is

slow enough for the run to last (at $\alpha = 0.30$ the two T_0 values were almost indistinguishable). The comparison with hill climbing also showed that the extra effort pays off exactly where it should: SA and hill climbing tie on the tiny instance, but SA's ability to accept uphill moves yields progressively better tours as the problem grows. In short, SA's superiority is not free — it is earned through a sensible choice of T_0, α and cooling schedule, and that choice, not the algorithm itself, is what determines whether the result is a near-optimal tour or a mediocre one.